Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:38 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 40 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

•[2]:

```
# ========================================================================
# Machine Learning Classification on the Adult (Census Income) Dataset
# Models : Logistic Regression, Decision Tree, KNN, Naive Bayes, SVM
# Metrics: Accuracy, AUC, Precision, Recall, F1, MCC
# Goal   : Predict whether a person earns >50K per year (1) or <=50K (0)
# ========================================================================

# ------------------------------------------------------------------------
#                       Import required libraries
# ------------------------------------------------------------------------
import pandas as pd                                  # data loading and tables
from sklearn.model_selection import train_test_split  # split data into train/test
from sklearn.compose import ColumnTransformer         # apply different preprocessing per column type
from sklearn.pipeline import Pipeline                 # chain preprocessing + model together
from sklearn.preprocessing import StandardScaler, OneHotEncoder

# The 5 classification models
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC
```

12:38 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:38 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 40 minutes ago

File Edit View Run Kernel Settings Help

Trusted

Code

JupyterLab | Python [conda env:base] *

```
from sklearn.naive_bayes import GaussianNB
from sklearn.svm import SVC

# Evaluation metrics
from sklearn.metrics import (accuracy_score, roc_auc_score, precision_score,
                             recall_score, f1_score, matthews_corrcoef)

# ------------------------------------------------------------------
#                       Load the dataset
# ------------------------------------------------------------------
df = pd.read_csv("adult_full.csv")   # the CSV created from the zip file
print("Original shape:", df.shape)

# ------------------------------------------------------------------
#                           Data cleaning
# ------------------------------------------------------------------
# Missing values in this dataset are written as '?', so convert them to
# proper missing values and remove those rows.
df = df.replace("?", pd.NA).dropna()
print("Shape after removing missing values:", df.shape)

# Convert the target column from text to numbers:
#   '>50K'  -> 1 (positive class)
#   '<=50K' -> 0 (negative class)
```

12:38 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:38 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 40 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
#   '>50K'  -> 1 (positive class)
#   '<=50K' -> 0 (negative class)
df["income"] = (df["income"] == ">50K").astype(int)

# ----------------------------------------------------------------------
#             Separate features (X) and target (y)
# ----------------------------------------------------------------------
X = df.drop(columns="income")   # all input columns
y = df["income"]                # the column we want to predict

# Identify numeric and categorical columns (they need different preprocessing)
num_cols = X.select_dtypes(include="number").columns.tolist()
cat_cols = X.select_dtypes(exclude="number").columns.tolist()
print("Numeric columns    :", num_cols)
print("Categorical columns:", cat_cols)

# ----------------------------------------------------------------------
#           Split into training and testing sets (80% / 20%)
# ----------------------------------------------------------------------
# stratify=y keeps the same >50K / <=50K ratio in both sets.
# random_state=42 makes the split reproducible.
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
print("Train size:", X_train.shape, "| Test size:", X_test.shape)
```

12:38 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:38 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 40 minutes ago

File Edit View Run Kernel Settings Help

Trusted

Code

JupyterLab | Python [conda env:base] *

```
    X, y, test_size=0.2, random_state=42, stratify=y
)
print("Train size:", X_train.shape, "| Test size:", X_test.shape)

# ---------------------------------------------------------------------
#                Build the preprocessing step
# ---------------------------------------------------------------------
# - Numeric columns    : standardize (mean 0, std 1) so large-valued columns
#                        (e.g. fnlwgt) do not dominate KNN / SVM / Logistic Regression
# - Categorical columns: one-hot encode into 0/1 columns
#                        (handle_unknown='ignore' avoids errors on unseen categories)
def make_preprocessor():
    return ColumnTransformer([
        ("num", StandardScaler(), num_cols),
        ("cat", OneHotEncoder(handle_unknown="ignore", sparse_output=False), cat_cols),
    ])

# ---------------------------------------------------------------------
#                  Define the 5 models
# ---------------------------------------------------------------------
models = {
    # 1. Linear model that predicts class probability using the sigmoid function
    "Logistic Regression": LogisticRegression(max_iter=1000),

    # 2. Tree of if/else rules; max_depth limits growth to prevent overfitting
```

12:38 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:38 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 40 minutes ago

File Edit View Run Kernel Settings Help

Trusted

Code

JupyterLab | Python [conda env:base] *

```
88
89      # 2. Tree of if/else rules; max_depth limits growth to prevent overfitting
90      "Decision Tree": DecisionTreeClassifier(max_depth=8, random_state=42),
91
92      # 3. Classifies a point by majority vote of its 7 nearest neighbours
93      "KNN": KNeighborsClassifier(n_neighbors=7),
94
95      # 4. Probabilistic model based on Bayes' theorem (assumes Gaussian features)
96      "Naive Bayes (Gaussian)": GaussianNB(),
97
98      # 5. Finds the best separating boundary; RBF kernel handles non-linear data
99      "SVM (RBF)": SVC(kernel="rbf", C=1.0),
100 }
101
102 # ----------------------------------------------------------------------
103 #           Train each model and calculate the evaluation metrics
104 # ----------------------------------------------------------------------
105 results = []   # will hold one row of metrics per model
106
107 for name, clf in models.items():
108
109     # Pipeline = preprocessing + model, so test data gets exactly the same
110     # transformation that was learned from the training data (no data leakage).
111     pipe = Pipeline([("pre", make_preprocessor()), ("clf", clf)])
112
```

12:38 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:38 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 40 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
    # 8b. Predict class labels (0 / 1) for the test set
    y_pred = pipe.predict(X_test)

    # 8c. Get continuous scores, which AUC needs.
    #     Most models give probabilities; SVC gives a decision-function score.
    if hasattr(pipe, "predict_proba"):
        y_score = pipe.predict_proba(X_test)[:, 1]   # probability of class 1
    else:
        y_score = pipe.decision_function(X_test)

    # 8d. Calculate all six metrics
    results.append({
        "Model":     name,
        "Accuracy":  accuracy_score(y_test, y_pred),    # % of correct predictions
        "AUC":       roc_auc_score(y_test, y_score),    # ability to rank positives above negatives
        "Precision": precision_score(y_test, y_pred),   # of predicted >50K, how many are truly >50K
        "Recall":    recall_score(y_test, y_pred),      # of actual >50K, how many were found
        "F1":        f1_score(y_test, y_pred),          # harmonic mean of precision and recall
        "MCC":       matthews_corrcoef(y_test, y_pred), # balanced score, good for imbalanced data
    })
    print(f"Finished: {name}")

# ------------------------------------------------------------------------
#             Show and save the comparison table
# ------------------------------
```

12:38 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:38 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 40 minutes ago

File Edit View Run Kernel Settings Help

Trusted

Code | JupyterLab | Python [conda env:base] *

```
# ----------------------------------------------------------------------
#             Show and save the comparison table
# ----------------------------------------------------------------------
results_df = pd.DataFrame(results).set_index("Model").round(4)
print("\n==================== Model Comparison ====================")
print(results_df.to_string())

results_df.to_csv("model_results.csv")
print("\nResults saved to model_results.csv")

# ----------------------------------------------------------------------
#        Identify the best model for each metric
# ----------------------------------------------------------------------
print("\nBest model per metric:")
for metric in results_df.columns:
    print(f"  {metric:<10}: {results_df[metric].idxmax()} ({results_df[metric].max():.4f})")

```

```
Original shape: (48842, 15)
Shape after removing missing values: (45222, 15)
Numeric columns    : ['age', 'fnlwgt', 'education_num', 'capital_gain', 'capital_loss', 'hours_per_week']
Categorical columns: ['workclass', 'education', 'marital_status', 'occupation', 'relationship', 'race', 'sex', 'native_country']
Train size: (36177, 14) | Test size: (9045, 14)
Finished: Logistic Regression
Finished: Decision Tree
Finished: KNN
```

12:38 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:38 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 41 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
Original shape: (48842, 15)
Shape after removing missing values: (45222, 15)
Numeric columns    : ['age', 'fnlwgt', 'education_num', 'capital_gain', 'capital_loss', 'hours_per_week']
Categorical columns: ['workclass', 'education', 'marital_status', 'occupation', 'relationship', 'race', 'sex', 'native_country']
Train size: (36177, 14) | Test size: (9045, 14)
Finished: Logistic Regression
Finished: Decision Tree
Finished: KNN
Finished: Naive Bayes (Gaussian)
Finished: SVM (RBF)

========== Model Comparison ==========
                        Accuracy     AUC  Precision  Recall      F1     MCC
Model
Logistic Regression       0.8450  0.9020     0.7344  0.5870  0.6525  0.5601
Decision Tree             0.8524  0.9004     0.7515  0.6044  0.6700  0.5820
KNN                       0.8302  0.8735     0.6781  0.5995  0.6364  0.5278
Naive Bayes (Gaussian)    0.6193  0.8389     0.3877  0.9242  0.5462  0.3891
SVM (RBF)                 0.8482  0.8957     0.7513  0.5794  0.6542  0.5668

Results saved to model_results.csv

Best model per metric:
  Accuracy  : Decision Tree (0.8524)
  AUC       : Logistic Regression (0.9020)
  Precision : Decision Tree (0.7515)
```

12:38 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:39 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 41 minutes ago

File Edit View Run Kernel Settings Help

Trusted

Code

JupyterLab | Python [conda env:base] *

```
Finished: Logistic Regression
Finished: Decision Tree
Finished: KNN
Finished: Naive Bayes (Gaussian)
Finished: SVM (RBF)

========== Model Comparison ==========
                        Accuracy     AUC  Precision  Recall      F1     MCC
Model
Logistic Regression       0.8450  0.9020     0.7344  0.5870  0.6525  0.5601
Decision Tree             0.8524  0.9004     0.7515  0.6044  0.6700  0.5820
KNN                       0.8302  0.8735     0.6781  0.5995  0.6364  0.5278
Naive Bayes (Gaussian)    0.6193  0.8389     0.3877  0.9242  0.5462  0.3891
SVM (RBF)                 0.8482  0.8957     0.7513  0.5794  0.6542  0.5668

Results saved to model_results.csv

Best model per metric:
  Accuracy  : Decision Tree (0.8524)
  AUC       : Logistic Regression (0.9020)
  Precision : Decision Tree (0.7515)
  Recall    : Naive Bayes (Gaussian) (0.9242)
  F1        : Decision Tree (0.6700)
  MCC       : Decision Tree (0.5820)
```

[4]:

```
1
2 # ----------------------------------------------------------------
```

12:39 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:39 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 41 minutes ago

File Edit View Run Kernel Settings Help

Trusted

Code

JupyterLab | Python [conda env:base] *

[4]:

```
# ------------------------------------------------------------------
#  OBSERVATIONS ON MODEL PERFORMANCE
# ------------------------------------------------------------------

observations = {
    "Logistic Regression": (
        "Logistic Regression provides a good baseline for income "
        "classification. It is computationally efficient and identifies "
        "relationships between demographic features and income. However, "
        "it may not capture complex non-linear relationships."
    ),

    "Decision Tree": (
        "Decision Tree identifies complex patterns and non-linear "
        "relationships between features. It is easy to interpret and "
        "understand. The maximum depth of 8 helps reduce overfitting."
    ),

    "KNN": (
        "KNN classifies individuals using their seven nearest neighbours. "
        "StandardScaler improves distance-based calculations. However, "
        "high-dimensional categorical features may affect its performance, "
        "and prediction can be computationally expensive."
    )
```

12:39 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:39 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 41 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
        "high-dimensional categorical features may affect its performance, "
        "and prediction can be computationally expensive."
    ),

    "Naive Bayes (Gaussian)": (
        "Gaussian Naive Bayes is fast and uses probability-based "
        "classification. However, its assumptions of conditional "
        "independence and Gaussian-distributed features may not accurately "
        "represent all features in the Adult dataset."
    ),

    "SVM (RBF)": (
        "SVM with an RBF kernel can identify complex non-linear decision "
        "boundaries. It can provide good classification performance "
        "but requires more computational time and memory for larger datasets."
    )
}

print("\n" + "=" * 80)
print("       OBSERVATIONS ON MODEL PERFORMANCE")
print("=" * 80)

for model_name, observation in observations.items():

    print("\nML Model Name:", model_name)
```

12:39 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:39 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 41 minutes ago

File Edit View Run Kernel Settings Help

Trusted

Code

JupyterLab | Python [conda env:base] *

```
46
47      print("\nML Model Name:", model_name)
48      print("Observation:", observation)
49
50      if model_name in results_df.index:
51
52          print("\nActual Evaluation Metrics:")
53
54          for metric in results_df.columns:
55              print(
56                  f"{metric:<12}: "
57                  f"{results_df.loc[model_name, metric]:.4f}"
58              )
59
60          print("-" * 80)
61
62
63  # ---------------------------------------------------------------------
64  #  IDENTIFY OVERALL WINNER
65  # ---------------------------------------------------------------------
66
67  print("\n" + "=" * 80)
68  print("                OVERALL MODEL WINNER")
69  print("=" * 80)
```

12:39 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:39 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 41 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
68 print("                OVERALL MODEL WINNER")
69 print("=" * 80)
70
71 # Calculate average of AUC, F1 and MCC
72 overall_scores = results_df[["AUC", "F1", "MCC"]].mean(axis=1)
73
74 # Identify winner
75 winner = overall_scores.idxmax()
76 winner_score = overall_scores.max()
77
78 print("\nModel Performance Based on AUC, F1 and MCC:")
79
80 print(overall_scores.round(4).to_string())
81
82 print("\nOverall Winner:", winner)
83 print("Average Score:", round(winner_score, 4))
84
85
86 # ------------------------------------------------------------------
87 #  GENERATE ASSIGNMENT CONCLUSION
88 # ------------------------------------------------------------------
89
90 print("\n" + "=" * 80)
91 print("                  ASSIGNMENT CONCLUSION")
92 print("=" * 80)
```

12:39 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:39 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 41 minutes ago

File Edit View Run Kernel Settings Help

Trusted

Code

JupyterLab | Python [conda env:base] *

```
print("\n" + "=" * 80)
print("                    ASSIGNMENT CONCLUSION")
print("=" * 80)

print("""
The Adult Census Income Dataset was used to implement five machine
learning classification algorithms.

All models were trained and evaluated using the same 80:20
training-testing split and preprocessing pipeline.

The models were compared using the following evaluation metrics:

1. Accuracy
2. AUC Score
3. Precision
4. Recall
5. F1 Score
6. Matthews Correlation Coefficient (MCC)
""")

print(
    f"The overall best-performing model is {winner}, "
    f"with an average AUC, F1 and MCC score of {winner_score:.4f}."
)
```

12:39 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3JIc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:39 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 41 minutes ago

File Edit View Run Kernel Settings Help

Trusted

Code

JupyterLab | Python [conda env:base] *

```
    f"with an average AUC, F1 and MCC score of {winner_score:.4f}."
)

print("""
The winning model was selected based on the average of AUC, F1
and MCC scores. This provides a more balanced evaluation than
accuracy alone, particularly for an imbalanced dataset.

Therefore, the selected model is considered the most suitable
among the five implemented classifiers for the Adult Census
Income Classification task.
""")

# ------------------------------------------------------------------
#  SAVE OBSERVATIONS AND CONCLUSION TO A TEXT FILE
# ------------------------------------------------------------------


# ------------------------------------------------------------------
#  SAVE OBSERVATIONS AND CONCLUSION TO A TEXT FILE
# ------------------------------------------------------------------

import os

# Save file in the current worki
```

12:39 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:39 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 41 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
141
142     file.write("OBSERVATIONS ON MODEL PERFORMANCE\n")
143     file.write("=" * 80 + "\n\n")
144
145     for model_name, observation in observations.items():
146
147         file.write(f"ML Model Name: {model_name}\n")
148         file.write(f"Observation: {observation}\n\n")
149
150         if model_name in results_df.index:
151
152             file.write("Evaluation Metrics:\n")
153
154             for metric in results_df.columns:
155
156                 file.write(
157                     f"{metric}: {results_df.loc[model_name, metric]:.4f}\n"
158                 )
159
160         file.write("\n" + "-" * 80 + "\n")
161
162     file.write("\nOVERALL WINNER\n")
163     file.write("=" * 80 + "\n")
164     file.write(f"Winner: {winner}\n")
165     file.write(f"Average Score: {winner_score:.4f}\n")
```

12:39 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:39 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 41 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
158                 )
159
160         file.write("\n" + "-" * 80 + "\n")
161
162     file.write("\nOVERALL WINNER\n")
163     file.write("=" * 80 + "\n")
164     file.write(f"Winner: {winner}\n")
165     file.write(f"Average Score: {winner_score:.4f}\n")
166
167 print("\nObservations saved successfully!")
168 print("File location:", os.path.abspath(file_path))
169
170
```

```
================================================================================
       OBSERVATIONS ON MODEL PERFORMANCE
================================================================================

ML Model Name: Logistic Regression
Observation: Logistic Regression provides a good baseline for income classification. It is computationally efficient and identifies relations
hips between demographic features and income. However, it may not capture complex non-linear relationships.

Actual Evaluation Metrics:
Accuracy    : 0.8450
```

12:39 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:39 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 42 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
169
170
```

```
================================================================================
       OBSERVATIONS ON MODEL PERFORMANCE
================================================================================

ML Model Name: Logistic Regression
Observation: Logistic Regression provides a good baseline for income classification. It is computationally efficient and identifies relations
hips between demographic features and income. However, it may not capture complex non-linear relationships.

Actual Evaluation Metrics:
Accuracy    : 0.8450
AUC         : 0.9020
Precision   : 0.7344
Recall      : 0.5870
F1          : 0.6525
MCC         : 0.5601
--------------------------------------------------------------------------------

ML Model Name: Decision Tree
Observation: Decision Tree identifies complex patterns and non-linear relationships between features. It is easy to interpret and understand.
The maximum depth of 8 helps reduce overfitting.

Actual Evaluation Metrics:
```

12:39 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:40 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 42 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
169
170
```

```
ML Model Name: Decision Tree
Observation: Decision Tree identifies complex patterns and non-linear relationships between features. It is easy to interpret and understand.
The maximum depth of 8 helps reduce overfitting.

Actual Evaluation Metrics:
Accuracy    : 0.8524
AUC         : 0.9004
Precision   : 0.7515
Recall      : 0.6044
F1          : 0.6700
MCC         : 0.5820
--------------------------------------------------------------------------------

ML Model Name: KNN
Observation: KNN classifies individuals using their seven nearest neighbours. StandardScaler improves distance-based calculations. However, h
igh-dimensional categorical features may affect its performance, and prediction can be computationally expensive.

Actual Evaluation Metrics:
Accuracy    : 0.8302
AUC         : 0.8735
Precision   : 0.6781
Recall      : 0.5995
F1          : 0.6364
```

12:40 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:40 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 42 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
169
170
```

```
Accuracy    : 0.8302
AUC         : 0.8735
Precision   : 0.6781
Recall      : 0.5995
F1          : 0.6364
MCC         : 0.5278
--------------------------------------------------------------------------------

ML Model Name: Naive Bayes (Gaussian)
Observation: Gaussian Naive Bayes is fast and uses probability-based classification. However, its assumptions of conditional independence and
Gaussian-distributed features may not accurately represent all features in the Adult dataset.

Actual Evaluation Metrics:
Accuracy    : 0.6193
AUC         : 0.8389
Precision   : 0.3877
Recall      : 0.9242
F1          : 0.5462
MCC         : 0.3891
--------------------------------------------------------------------------------

ML Model Name: SVM (RBF)
Observation: SVM with an RBF kernel can identify complex non-linear decision boundaries. It can provide good classification performance but r
```

12:40 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:40 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 42 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
169
170
```

```
MCC         : 0.3891
--------------------------------------------------------------------------------

ML Model Name: SVM (RBF)
Observation: SVM with an RBF kernel can identify complex non-linear decision boundaries. It can provide good classification performance but r
equires more computational time and memory for larger datasets.

Actual Evaluation Metrics:
Accuracy    : 0.8482
AUC         : 0.8957
Precision   : 0.7513
Recall      : 0.5794
F1          : 0.6542
MCC         : 0.5668
--------------------------------------------------------------------------------

================================================================================
                OVERALL MODEL WINNER
================================================================================

Model Performance Based on AUC, F1 and MCC:
Model
Logistic Regression       0.7049
```

12:40 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:40 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 42 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
169
170
```

```
F1          : 0.6542
MCC         : 0.5668
--------------------------------------------------------------------------------

================================================================================
                OVERALL MODEL WINNER
================================================================================

Model Performance Based on AUC, F1 and MCC:
Model
Logistic Regression       0.7049
Decision Tree             0.7175
KNN                       0.6792
Naive Bayes (Gaussian)    0.5914
SVM (RBF)                 0.7056

Overall Winner: Decision Tree
Average Score: 0.7175

================================================================================
                ASSIGNMENT CONCLUSION
================================================================================
```

12:40 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:40 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 42 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

```
169
170
```

```
================================================================================
                  ASSIGNMENT CONCLUSION
================================================================================

The Adult Census Income Dataset was used to implement five machine
learning classification algorithms.

All models were trained and evaluated using the same 80:20
training-testing split and preprocessing pipeline.

The models were compared using the following evaluation metrics:

1. Accuracy
2. AUC Score
3. Precision
4. Recall
5. F1 Score
6. Matthews Correlation Coefficient (MCC)

The overall best-performing model is Decision Tree, with an average AUC, F1 and MCC score of 0.7175.

The winning model was selected based on the average of AUC, F1
```

12:40 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:40 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 42 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

The Adult Census Income dataset was used to implement five machine learning classification algorithms.

All models were trained and evaluated using the same 80:20 training-testing split and preprocessing pipeline.

The models were compared using the following evaluation metrics:

1. Accuracy
2. AUC Score
3. Precision
4. Recall
5. F1 Score
6. Matthews Correlation Coefficient (MCC)

The overall best-performing model is Decision Tree, with an average AUC, F1 and MCC score of 0.7175.

The winning model was selected based on the average of AUC, F1 and MCC scores. This provides a more balanced evaluation than accuracy alone, particularly for an imbalanced dataset.

Therefore, the selected model is considered the most suitable among the five implemented classifiers for the Adult Census Income Classification task.

Observations saved successfully!

12:40 PM 11/5/2026

Lab - 8310 - s2026aj05499 - 55497 - Work - Microsoft Edge

https://g1.prayogshala.bits-pilani.ac.in/web-rdp/#/client/ODAxOTYAYwBwb3N0Z3Jlc3Fs?token=692E6F7F2C31465367ECEF93BA4561583EEE5F278E1DEAB0D7BF254AC7452BC0

Applications | ML_Assignment2 — Mo... | Downloads - Thunar | 2026-11-05 12:40 | s2026aj05499

Home | ML_Assignment2 | Jupyter Notebook - Term

localhost:8888/notebooks/Ass-2ML/ML_Assignment2.ipynb?

jupyter ML_Assignment2 Last Checkpoint: 42 minutes ago

File Edit View Run Kernel Settings Help Trusted

Code JupyterLab Python [conda env:base] *

The models were compared using the following evaluation metrics:

1. Accuracy
2. AUC Score
3. Precision
4. Recall
5. F1 Score
6. Matthews Correlation Coefficient (MCC)

The overall best-performing model is Decision Tree, with an average AUC, F1 and MCC score of 0.7175.

The winning model was selected based on the average of AUC, F1
and MCC scores. This provides a more balanced evaluation than
accuracy alone, particularly for an imbalanced dataset.

Therefore, the selected model is considered the most suitable
among the five implemented classifiers for the Adult Census
Income Classification task.

Observations saved successfully!
File location: /home/s2026aj05499/Ass-2ML/model_observations.txt

12:40 PM 11/5/2026